


# Data Visualization I

Michael C. Hackett


Assistant Professor, Computer Science

Community  
College  
*of* Philadelphia




# ggplot2

- ggplot2 is a data visualization package in R
- Allows for declaratively creating graphics
  - Based on the text [The Grammar of Graphics](#)
- *“You provide the data, tell ggplot2 how to map variables to aesthetics, what graphical primitives to use, and it takes care of the details.”*
  - Project homepage: <https://ggplot2.tidyverse.org/>




# ggplot2

- ggplot2 is installed along with the tidyverse:  
`install.packages("tidyverse")`
- Can be installed as a stand-alone package:  
`install.packages("ggplot2")`
- Extensions:  
<https://exts.ggplot2.tidyverse.org/gallery/>



# ggplot2


- ggplot2 is loaded along with the rest of the tidyverse:  
**library(tidyverse)**
- Can be loaded by itself:  
**library(ggplot2)**
- ggplot2 has a sample data frame for demonstration purposes
  - The **mpg** dataset contains observations collected by the US Environmental Protection Agency on 38 models of cars



# ggplot2


- If tidyverse was loaded:  
`library(tidyverse)`  
`ggplot2::mpg`
- If ggplot2 was loaded by itself:  
`library(ggplot2)`  
`mpg`
- We'll assume ggplot2 was loaded by itself for the remainder of the lecture

# ggplot2




```
> library(ggplot2)
Use suppressPackageStartupMessages() to eliminate package startup messages
Warning message:
package 'ggplot2' was built under R version 3.6.3
> mpg
# A tibble: 234 x 11
  manufacturer model      displ  year   cyl trans      drv   cty   hwy fl  class
  <chr>         <chr>    <dbl> <int> <int> <chr>    <chr> <int> <int> <chr> <chr>
1 audi         a4          1.8  1999     4 auto(l5)  f     18    29 p    compact
2 audi         a4          1.8  1999     4 manual(m5) f     21    29 p    compact
3 audi         a4          2    2008     4 manual(m6) f     20    31 p    compact
4 audi         a4          2    2008     4 auto(av)   f     21    30 p    compact
5 audi         a4          2.8  1999     6 auto(l5)  f     16    26 p    compact
6 audi         a4          2.8  1999     6 manual(m5) f     18    26 p    compact
7 audi         a4          3.1  2008     6 auto(av)   f     18    27 p    compact
8 audi         a4 quattro    1.8  1999     4 manual(m5) 4     18    26 p    compact
9 audi         a4 quattro    1.8  1999     4 auto(l5)   4     16    25 p    compact
10 audi         a4 quattro     2    2008     4 manual(m6) 4     20    28 p    compact
# ... with 224 more rows
```

- The **displ** column is the engine size (in liters)
- The **hwy** column is the highway gas milage in miles-per-gallon



# ggplot2

- We begin creating a plot with the **ggplot()** function
  - This creates a coordinate system that layers can be added on to
- The first argument to the ggplot function is the data we wish to plot  
**ggplot(data = mpg)**
- Now that the plot has its data, layers are added that specify how to data is to be displayed.
  - Layers of data are referred to as *geometries* or ***geoms***




# ggplot2

- ggplot2's **geom\_point()** function adds a layer of points to a plot
  - Effectively creating a scatterplot
- The first argument to the **geom\_point** function is the mapping
  - Will define how variables are mapped to different aesthetics or visual properties on this layer

**geom\_point(mapping = aes())**

- The **aes** function shown here specifies the aesthetics of the layer






# ggplot2 - Aesthetics

- The x and y arguments to the aes function specify which columns to use for the x and y axes of our scatterplot

```
geom_point(mapping = aes(x = displ , y = hwy))
```



f + geom\_point()


x, y, alpha, color, fill, shape, size

Aesthetics for geom\_point from  
cheat sheet


- Now, we simply add this layer to the plot:


```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ , y = hwy))
```

# ggplot2 - Aesthetics



```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ , y = hwy))
```







# ggplot2 - Aesthetics

- The color argument to the `geom_point` function will specify the color of all the data points

```
geom_point(mapping = aes(x = displ , y = hwy), color = "red")
```

```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ , y = hwy), color = "red")
```






# ggplot2 - Aesthetics

- The color argument to the *aes function* will specify how to color the individual data points


```
geom_point(mapping = aes(x = displ , y = hwy, color = class))
```


- This will apply a unique color to each point, based on the class column of each observation

# ggplot2 - Aesthetics



```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ , y = hwy, color = class))
```






# ggplot2 - Aesthetics

- The alpha argument to the aes function will specify the opacity of the data points


```
geom_point(mapping = aes(x = displ , y = hwy, alpha = class))
```


- This will apply an opacity to each point, based on the class column of each observation

# ggplot2 - Aesthetics



```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ , y = hwy, alpha = class))
```





# ggplot2 - Aesthetics

- The size argument to the aes function will specify the size of the data points

```
geom_point(mapping = aes(x = displ , y = hwy, size = class))
```


- This will apply a size to each point, based on the class column of each observation




# ggplot2 - Aesthetics



```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ , y = hwy, size = class))
```






# ggplot2 - Aesthetics

- The shape argument to the aes function will specify the shapes of the data points


```
geom_point(mapping = aes(x = displ , y = hwy, shape = class))
```

- This will apply a shape to each point, based on the class column of each observation


# ggplot2 - Aesthetics



```
ggplot(data = mpg) + geom_point(mapping = aes(x = displ , y = hwy, shape = class))
```




- Note: Only displays up to 6 shapes (hence why SUV is not shown)




# ggplot2 - Geoms

- The `geom_smooth` function plots a layer of a smooth, fitted line

```
ggplot(data = mpg) + geom_smooth(mapping = aes(x = displ , y = hwy))
```







# ggplot2 - Geoms

- The layer can be plotted on top of other layers

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ , y = hwy)) +  
  geom_smooth(mapping = aes(x = displ , y = hwy))
```







# ggplot2 - Geoms

- There is some redundancy in the last slide
  - Both layers needed to be told the x and y values.
- Instead, the mapping can be moved into the original ggplot
  - Now, the data and a default mapping/aesthetic is passed up to any added layers

```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point() +  
  geom_smooth()
```





(Same plot)



# ggplot2 - Geoms

```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point(mapping = aes(color = class)) +  
  geom_smooth()
```






# ggplot2 - Labels

- Labels such as title, subtitle, captions, and axis labels can be added using ggplot2's lab function

```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point() +  
  geom_smooth() +  
  labs(title = "Fuel efficiency generally decreases with engine size")
```






# ggplot2 - Labels





```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) + geom_point() + geom_smooth() +  
labs(title = "Fuel efficiency generally decreases with engine size",
```


```
  subtitle = "Two seaters (sports cars) are an exception because of their light weight",  
  caption = "Data from fueleconomy.gov")
```



(Shown larger on next slide)

# ggplot2 - Labels








# ggplot2 - Labels

```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) + geom_point() + geom_smooth() +  
labs(title = "Fuel efficiency generally decreases with engine size",  
      subtitle = "Two seaters (sports cars) are an exception because of their light weight",  
      caption = "Data from fueleconomy.gov",  
      x = "Engine Displacement (L)", y = "Highway Fuel Economy (mpg)")
```

(Shown on next slide)


# ggplot2 - Labels





# ggplot2 - Scales

- Scales in ggplot2 control how a plot maps data values to the visual values of an aesthetic.
- The general format of a scale function is  
`scale*_function()`
- For example, `scale_x_reverse` will reverse the x axis





# ggplot2 - Scales


```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point() +  
  geom_smooth() +  
  labs(title = "Fuel efficiency generally decreases with engine size",  
        subtitle = "Two seaters (sports cars) are an exception because of their light weight",  
        caption = "Data from fueleconomy.gov",  
        x = "Engine Displacement (L)", y = "Highway Fuel Economy (mpg)") +  
  scale_x_reverse()
```

(Shown on next slide)

# ggplot2 - Scales




(Not really helpful for this graph, but it demonstrates a scale in ggplot2)



# ggplot2 - Scales

- A scale can be used to adjust the tick marks for the scatterplot.
- The **scale\_y\_continuous** function will allow us to change the breaks and limits for the y axis
  - **scale\_x\_continuous** would do the same for the x axis
  - Use **scale\_x\_discrete** and **scale\_y\_discrete** for axes that use discrete variables (like text)
- The **breaks** argument is a vector that specifies the labels
- The **limits** argument is a vector that specifies the axis' range.








# ggplot2 - Scales

```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point() +  
  geom_smooth() +  
  labs(title = "Fuel efficiency generally decreases with engine size",  
        subtitle = "Two seaters (sports cars) are an exception because of their light weight",  
        caption = "Data from fueleconomy.gov",  
        x = "Engine Displacement (L)", y = "Highway Fuel Economy (mpg)") +  
  scale_y_continuous(breaks=seq(0, 50, 5), limits=c(5,50))
```

(Shown on next slide)

# ggplot2 - Scales






# ggplot2 - Scales

```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point() +  
  geom_smooth() +  
  labs(title = "Fuel efficiency generally decreases with engine size",  
        subtitle = "Two seaters (sports cars) are an exception because of their light weight",  
        caption = "Data from fueleconomy.gov",  
        x = "Engine Displacement (L)", y = "Highway Fuel Economy (mpg)") +  
  scale_y_continuous(breaks=seq(0, 50, 5), limits=c(5,50)) +  
  scale_x_continuous(breaks=seq(1.5, 7.1, 0.5), limits=c(1.5, 7))
```


(Shown on next slide)


# ggplot2 - Scales



Fuel efficiency generally decreases with engine size


Two seaters (sports cars) are an exception because of their light weight






# ggplot2 - Legends

- There are a number of ways to control how the legend is displayed.







# ggplot2 - Legends


- Changing the Legend title:

```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point() +  
  geom_smooth() +  
  labs(x = "Engine Displacement (L)", y = "Highway Fuel Economy (mpg)",  
        color = "Car Class") +  
  scale_y_continuous(breaks=seq(0, 50, 5), limits=c(5,50)) +  
  scale_x_continuous(breaks=seq(1.5, 7.1, 0.5), limits=c(1.5, 7))
```

(Shown on next slide)

# ggplot2 - Legends





# ggplot2 - Legends



- Changing the Legend position:
  - “bottom”, “top”, “left”, or “right”


```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point() +  
  geom_smooth() +  
  labs(x = "Engine Displacement (L)", y = "Highway Fuel Economy (mpg)",  
        color = "Car Class") +  
  scale_y_continuous(breaks=seq(0, 50, 5), limits=c(5,50)) +  
  scale_x_continuous(breaks=seq(1.5, 7.1, 0.5), limits=c(1.5, 7)) +  
  theme(legend.position = "bottom")
```

(Shown on next slide)



# ggplot2 - Legends





# ggplot2 - Legends

- We can change the labels in the legend, but it will also require specifying the colors of each:

```
ggplot(data = mpg, mapping = aes(x = displ , y = hwy)) +  
  geom_point() +  
  geom_smooth() +  
  labs(x = "Engine Displacement (L)", y = "Highway Fuel Economy (mpg)",  
        color = "Car Class") +  
  scale_y_continuous(breaks=seq(0, 50, 5), limits=c(5,50)) +  
  scale_x_continuous(breaks=seq(1.5, 7.1, 0.5), limits=c(1.5, 7)) +  
  theme(legend.position = "bottom") +  
  scale_color_manual(labels=c("Two Seater", "Compact", "Mid-Size",  
                             "Minivan", "Pickup", "Sub-Compact", "SUV"),  
                    values=c("red", "orange", "yellow", "green",  
                             "blue", "purple", "violet"))
```

(Shown on next slide)

# ggplot2 - Legends

